

#libraries

```
import pandas as pd
import numpy as npp
import matplotlib.pyplot as plt
%matplotlib inline
```

#load dataset

```
df = pd.read_csv(r"/content/Churn_Modelling.csv")
```

```
df.head(10)
```

RowNumber CustomerId Surname CreditScore Geography Gender Age

```
\
0      1  15634602  Hargrave      619  France Female 42
1      2  15647311   Hill      608  Spain Female 41
2      3  15619304   Onio      502  France Female 42
3      4  15701354   Boni      699  France Female 39
4      5  15737888  Mitchell      850  Spain Female 43
5      6  15574012   Chu      645  Spain  Male 44
6      7  15592531  Bartlett      822  France  Male 50
7      8  15656148  Obinna      376  Germany Female 29
8      9  15792365   He      501  France  Male 44
9     10  15592389   H?      684  France  Male 27
```

```
Tenure  Balance NumOfProducts  HasCrCard  IsActiveMember \
0      2      0.00           1           1           1
1      1  83807.86           1           0           1
2      8 159660.80           3           1           0
3      1      0.00           2           0           0
4      2 125510.82           1           1           1
5      8 113755.78           2           1           0
6      7      0.00           2           1           1
7      4 115046.74           4           1           0
8      4 142051.07           2           0           1
9      2 134603.88           1           1           1
```

EstimatedSalary Exited

```
0      101348.88      1
1      112542.58      0
2      113931.57      1
```

```
3      93826.63      0
4      79084.10      0
5      149756.71      1
6      10062.80      0
7      119346.88      1
8      74940.50      0
9      71725.73      0
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	RowNumber	10000 non-null	int64
1	CustomerId	10000 non-null	int64
2	Surname	10000 non-null	object
3	CreditScore	10000 non-null	int64
4	Geography	10000 non-null	object
5	Gender	10000 non-null	object
6	Age	10000 non-null	int64
7	Tenure	10000 non-null	int64
8	Balance	10000 non-null	float64
9	NumOfProducts	10000 non-null	int64
10	HasCrCard	10000 non-null	int64
11	IsActiveMember	10000 non-null	int64
12	EstimatedSalary	10000 non-null	float64
13	Exited	10000 non-null	int64

```
dtypes: float64(2), int64(9), object(3)
```

```
memory usage: 1.1+ MB
```

```
#Visualizations
```


```
#Univariate Analysis
```

```
import seaborn as sns
```


```
sns.kdeplot(df['CreditScore'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fc4a0cd2790>
```

```
#Bi - Variate Analysis
plt.bar(df.CustomerId, df.CreditScore)
plt.title('CreditScore')
plt.xlabel('CustomerId')
plt.ylabel('CreditScore')
Text(0, 0.5, 'CreditScore')
```




```
sns.lmplot(x='Tenure', y='Balance', data=df, hue='Exited', size=8)
/usr/local/lib/python3.7/dist-packages/seaborn/regression.py:581:
UserWarning: The `size` parameter has been renamed to `height`; please
update your code.
  warnings.warn(msg, UserWarning)
<seaborn.axisgrid.FacetGrid at 0x7fc4a149e2d0>
```



```
#Multi - Variate Analysis
```

```
ax =  
df[["CreditScore","Age","Tenure","Balance"]].plot(figsize=(80,40))  
ax.legend(loc='center left', bbox_to_anchor=(1, 0.5));
```



```
df.isnull().sum()
RowNumber      0
CustomerId     0
Surname        0
CreditScore    0
Geography      0
Gender         0
Age           0
Tenure        0
Balance       0
NumOfProducts 0
HasCrCard     0
IsActiveMember 0
EstimatedSalary 0
Exited        0
dtype: int64
plt.figure(figsize=(15,13))
sns.heatmap(df.corr(),annot=True,cmap='BuPu')
plt.show()
```

[illegible]

```

0      0      1      112542.58      0
2      1      0      113931.57      1
3      0      0      93826.63      0
4      1      1      79084.10      0
df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   CreditScore  10000 non-null  int64
1   Geography    10000 non-null  object
2   Gender        10000 non-null  object
3   Age          10000 non-null  int64
4   Tenure       10000 non-null  int64
5   Balance      10000 non-null  float64
6   NumOfProducts 10000 non-null  int64
7   HasCrCard    10000 non-null  int64
8   IsActiveMember 10000 non-null  int64
9   EstimatedSalary 10000 non-null float64
10  Exited       10000 non-null  int64
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
df["Geography"].unique()
array(['France', 'Spain', 'Germany'], dtype=object)

df["Gender"].unique()
array(['Female', 'Male'], dtype=object)

geo=pd.get_dummies(df["Geography"],drop_first=False)

geo.head()
   France  Germany  Spain
0      1      0      0
1      0      0      1
2      1      0      0
3      1      0      0
4      0      0      1

gen=pd.get_dummies(df["Gender"],drop_first=False)
df=pd.concat([df, geo,gen], axis=1)

df
   CreditScore  Geography  Gender  Age  Tenure  Balance
NumOfProducts \

```


1	608	Spain	Female	41	1	83807.86	
1							
2	502	France	Female	42	8	159660.80	
3							
3	699	France	Female	39	1	0.00	
2							
4	850	Spain	Female	43	2	125510.82	
1							
...	
...							
9995	771	France	Male	39	5	0.00	
2							
9996	516	France	Male	35	10	57369.61	
1							
9997	709	France	Female	36	7	0.00	
1							
9998	772	Germany	Male	42	3	75075.31	
2							
9999	792	France	Female	28	4	130142.79	
1							
HasCrCard IsActiveMember EstimatedSalary Exited France							
Germany \							
0	1		1		101348.88	1	1
0							
1	0		1		112542.58	0	0
0							
2	1		0		113931.57	1	1
0							
3	0		0		93826.63	0	1
0							
4	1		1		79084.10	0	0
0							
...
...							
9995	1		0		96270.64	0	1
0							
9996	1		1		101699.77	0	1
0							
9997	0		1		42085.58	1	1
0							
9998	1		0		92888.52	1	0
1							
9999	1		0		38190.78	0	1
0							
Spain Female Male							
0	0	1	0				

```
1 1 1 0
2 0 1 0
3 0 1 0
4 1 1 0
```

```
... ..
9995 0 0 1
9996 0 0 1
9997 0 1 0
9998 0 0 1
9999 0 1 0
```

```
[10000 rows x 16 columns]
```

```
df.drop(["Geography","Gender"], axis=1, inplace=True)
```

```
df.head()
```

```
CreditScore Age Tenure Balance NumOfProducts HasCrCard \
0 619 42 2 0.00 1 1
1 608 41 1 83807.86 1 0
2 502 42 8 159660.80 3 1
3 699 39 1 0.00 2 0
4 850 43 2 125510.82 1 1
IsActiveMember EstimatedSalary Exited France Germany Spain
```

```
Female \
0 1 101348.88 1 1 0 0
1
1 1 112542.58 0 0 0 1
1
2 0 113931.57 1 1 0 0
1
3 0 93826.63 0 1 0 0
1
4 1 79084.10 0 0 0 1
```

```
Male
0 0
1 0
2 0
3 0
4 0
```

```
x=df.drop('Exited',axis=1)
```

```
x
```

```
CreditScore Age Tenure Balance NumOfProducts HasCrCard \
0 619 42 2 0.00 1 1
1 608 41 1 83807.86 1 0
```

3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	
...		
9995	771	39	5	0.00	2	1	
9996	516	35	10	57369.61	1	1	
9997	709	36	7	0.00	1	0	
9998	772	42	3	75075.31	2	1	
9999	792	28	4	130142.79	1	1	
IsActiveMember EstimatedSalary France Germany Spain Female							
Male							
0	1	101348.88	1	0	0	1	
0							
1	1	112542.58	0	0	1	1	
0							
2	0	113931.57	1	0	0	1	
0							
3	0	93826.63	1	0	0	1	
0							
4	1	79084.10	0	0	1	1	
0							
...		
...							
9995	0	96270.64	1	0	0	0	
1							
9996	1	101699.77	1	0	0	0	
1							
9997	1	42085.58	1	0	0	1	
0							
9998	0	92888.52	0	1	0	0	
[10000 rows x 13 columns]							
9999	0	38190.78	1	0	0	1	
0							
Y							
0	1						
1	0						
2	1						
3	0						
4	0						
..							
9995	0						
9996	0						
9997	1						
9998	1						

```
9999 0
Name:Exited, Length: 10000, dtype: int64
df.shape
(10000, 14)
x.shape
(10000, 13)
y.shape
(10000,)
from sklearn.model_selection import train_test_split
x_train,x_test, y_train,y_test = train_test_split(x,y,
test_size=0.2,random_state=0)
x_train.shape
(8000, 13)
x_test.shape
(2000, 13)
y_test.shape
(2000,)
from sklearn.preprocessing import StandardScaler

sc = StandardScaler()
x_train = sc.fit_transform(x_train)

x_train
array([[ 0.16958176, -0.46460796, 0.00666099, ..., 1.74309049,
        1.09168714, -1.09168714],
       [-2.30455945, 0.30102557, -1.37744033, ..., -0.57369368,
        -0.91601335, 0.91601335],
       [-1.19119591, -0.94312892, -1.031415 , ..., -0.57369368,
        1.09168714, -1.09168714],
       ...,
       [ 0.9015152 , -0.36890377, 0.00666099, ..., -0.57369368,
        -0.91601335, 0.91601335],
       [-0.62420521, -0.08179119, 1.39076231, ..., 1.74309049,
        1.09168714, -1.09168714],
       [-0.28401079, 0.87525072, -1.37744033, ..., -0.57369368,
        1.09168714, -1.09168714]])
x_test = sc.transform(x_test)
```

```
x_test
array([[ -0.55204276, -0.36890377,  1.04473698, ..., -0.57369368,
         1.09168714, -1.09168714],
       [-1.31490297,  0.10961719, -1.031415 , ..., -0.57369368,
         1.09168714, -1.09168714],
       [ 0.57162971,  0.30102557,  1.04473698, ...,  1.74309049,
         1.09168714, -1.09168714],
       ...,
       [-0.74791227, -0.27319958, -1.37744033, ...,  1.74309049,
        -0.91601335,  0.91601335],
       [-0.00566991, -0.46460796, -0.33936434, ..., -0.57369368,
        -0.91601335,  0.91601335],
       [-0.79945688, -0.84742473,  1.04473698, ..., -0.57369368,
        -0.91601335,  0.91601335]])
```